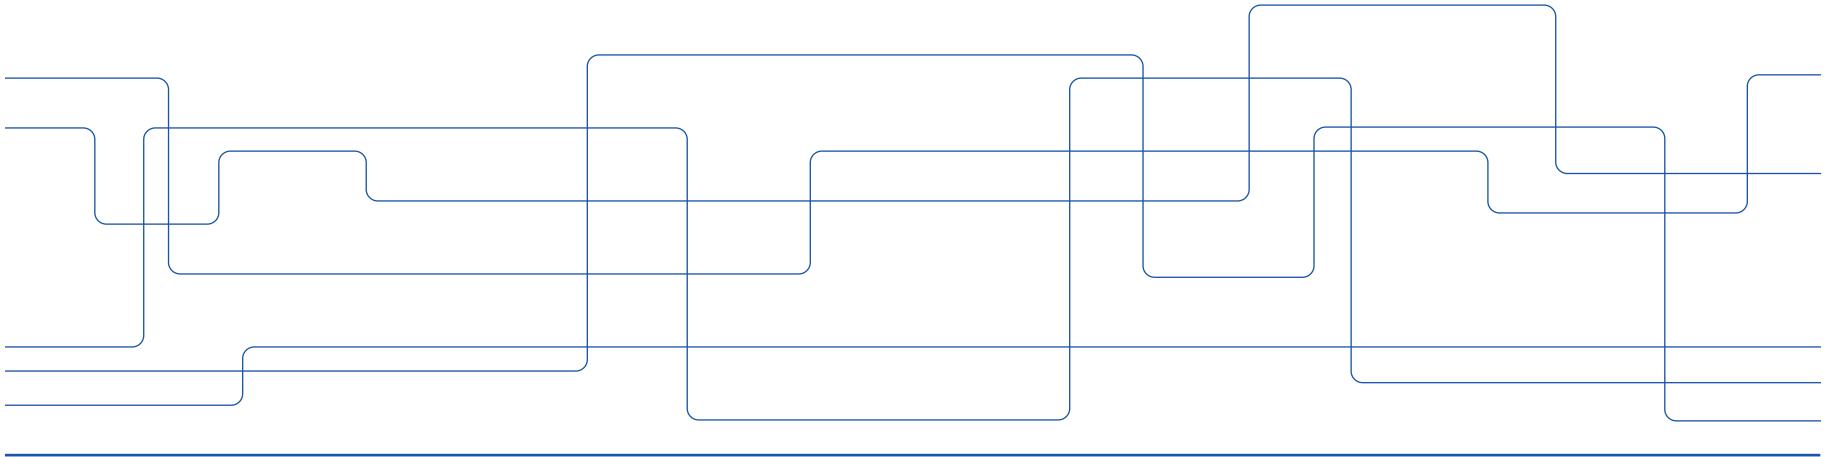




DD2358 – Using the cProfile Module

Stefano Markidis

KTH Royal Institute of Technology





Intended Learning Outcomes

- To obtain profiling information about the most used functions in the code with the `cProfile` Python tool
- To analyze the `cProfile` output and statistics and understand the impact of different function on the overall application timing
- Use the `SnakeViz` tool to visualize the results of `cProfile`



The cProfile tool

- cProfile is a built-in profiling tool in the standard library.
 - It hooks into the python interpreter in CPython to measure the time taken to run every function that it sees.
- profile is the original and slower pure Python profiler
- cProfile has the same interface as profile and is **written in C for a lower overhead.**
- If you're curious about the history of these libraries, see [Armin Rigo's 2005 request](#) to include cProfile in the standard library.

[Python-Dev] s/hotshot/lspref

Armin Rigo [arigo at tunes.org](mailto:arigo@tunes.org)

Sat Nov 19 19:08:55 CET 2005

- Previous message: [\[Python-Dev\] How to stay almost backwards compatible with a](#)
- Next message: [\[Python-Dev\] s/hotshot/lspref](#)
- Messages sorted by: [\[date\]](#) [\[thread\]](#) [\[subject\]](#) [\[author\]](#)

Hi!

The current Python profilers situation is a mess.

'profile.Profile' is the ages-old pure Python profiler. At the end of a run, it builds a dict that is inspected by 'pstats.Stats'. It has some recent support for profiling C calls, which however make it crash in some cases [1]. And of course it's slow (makes a run take about 10x longer).

'hotshot', new from 2.2, is quite faster (reportedly, only 30% added overhead). The log file is then loaded and turned into an instance of



Hypothesis when Profiling

- A good practice when profiling is to generate a ***hypothesis*** about the speed of parts of your code before you profile it.
- Let's hypothesize that `calculate_z_serial_purepython` is the slowest part of the code.
 - In that function, we do a lot of **dereferencing** and make **many calls to basic arithmetic operators** and the **`abs`** function.
 - > *These will probably show up as consumers of CPU resources.*
- Here, we'll use the `cProfile` module to run the code.
 - The output is spartan but helps us figure out where to analyze further.

```
def calculate_z_serial_purepython(maxiter, zs, cs):  
    """Calculate output list using Julia update rule"""  
    output = [0] * len(zs)  
    for i in range(len(zs)):  
        n = 0  
        z = zs[i]  
        c = cs[i]  
        while abs(z) < 2 and n < maxiter:  
            z = z * z + c  
            n += 1  
        output[i] = n  
    return output
```

```
python -m cProfile -s cumulative JuliaSet.py
```

The -s cumulative flag tells `cProfile` to **sort by cumulative time spent** inside each function; this gives us a view into the slowest parts of a section of code. The `cProfile` output is written to screen directly after our usual print results

```
stef@Stefs-MacBook-Air Codes % python -m cProfile -s cumulative JuliaSet.py
```

```
Length of x: 1000
```

```
Total elements: 1000000
```

```
calculate_z_serial_purepython took 4.945113182067871 seconds
```

```
36221995 function calls in 5.263 seconds
```

Ordered by: cumulative time

ncalls	tottime	percall	cumtime	percall	filename:lineno(function)
1	0.000	0.000	5.263	5.263	{built-in method builtins.exec}
1	0.013	0.013	5.263	5.263	JuliaSet.py:1(<module>)
1	0.252	0.252	5.249	5.249	JuliaSet.py:21(calc_pure_python)
1	3.975	3.975	4.945	4.945	JuliaSet.py:59(calculate_z_serial_purepython)
34219980	0.970	0.000	0.970	0.000	{built-in method builtins.abs}
2002000	0.048	0.000	0.048	0.000	{method 'append' of 'list' objects}
1	0.004	0.004	0.004	0.004	{built-in method builtins.sum}
3	0.000	0.000	0.000	0.000	{built-in method builtins.print}
2	0.000	0.000	0.000	0.000	{built-in method time.time}
4	0.000	0.000	0.000	0.000	{built-in method builtins.len}
1	0.000	0.000	0.000	0.000	{method 'disable' of '_lsprof.Profiler' objects}


```

stef@Stefs-MacBook-Air Codes % python -m cProfile -s cumulative JuliaSet.py
Length of x: 1000
Total elements: 1000000
calculate_z_serial_purepython took 4.945113182067871 seconds
36221995 function calls in 5.263 seconds

```

Ordered by: cumulative time

ncalls	tottime	percall	cumtime	percall	filename:lineno(function)
1	0.000	0.000	5.263	5.263	{built-in method builtins.exec}
1	0.013	0.013	5.263	5.263	JuliaSet.py:1(<module>)
1	0.252	0.252	5.249	5.249	JuliaSet.py:21(calc_pure_python)
1	3.975	3.975	4.945	4.945	JuliaSet.py:59(calculate_z_serial_purepython)
34219980	0.970	0.000	0.970	0.000	{built-in method builtins.abs}
2002000	0.048	0.000	0.048	0.000	{method 'append' of 'list' objects}
1	0.004	0.004	0.004	0.004	{built-in method builtins.sum}
3	0.000	0.000	0.000	0.000	{built-in method builtins.print}
2	0.000	0.000	0.000	0.000	{built-in method time.time}
4	0.000	0.000	0.000	0.000	{built-in method builtins.len}
1	0.000	0.000	0.000	0.000	{method 'disable' of '_lsprof.Profiler' objects}




- We can see that the entry point to the code `JuliasSet.py` on line 1 takes a total of 5.2 seconds.
- This is just the `__main__` call to `calc_pure_python`.
- `ncalls` is 1, indicating that this line is executed only once.

```
stef@Stefs-MacBook-Air Codes % python -m cProfile -s cumulative JuliaSet.py
Length of x: 1000
Total elements: 1000000
calculate_z_serial_purepython took 4.945113182067871 seconds
36221995 function calls in 5.263 seconds
```

Ordered by: cumulative time

ncalls	tottime	percall	cumtime	percall	filename:lineno(function)
1	0.000	0.000	5.263	5.263	{built-in method builtins.exec}
1	0.013	0.013	5.263	5.263	JuliaSet.py:1(<module>)
1	0.252	0.252	5.249	5.249	JuliaSet.py:21(calc_pure_python)
1	3.975	3.975	4.945	4.945	JuliaSet.py:59(calculate_z_serial_purepython)
34219980	0.970	0.000	0.970	0.000	{built-in method builtins.abs}
2002000	0.048	0.000	0.048	0.000	{method 'append' of 'list' objects}
1	0.004	0.004	0.004	0.004	{built-in method builtins.sum}
3	0.000	0.000	0.000	0.000	{built-in method builtins.print}
2	0.000	0.000	0.000	0.000	{built-in method time.time}
4	0.000	0.000	0.000	0.000	{built-in method builtins.len}
1	0.000	0.000	0.000	0.000	{method 'disable' of '_lsprof.Profiler' objects}



- Inside `calc_pure_python`, the call to `calculate_z_serial_purepython` consumes 4.95 seconds.
 - > *Both functions are called only once.*
- We can derive that approximately 0.3 second is spent on lines of code inside `calc_pure_python`, separate to calling the CPU-intensive `calculate_z_serial_purepython` function.
 - However, we can't derive *which* lines take the time inside the function using `cProfile`.

```
stef@Stefs-MacBook-Air Codes % python -m cProfile -s cumulative JuliaSet.py
```

```
Length of x: 1000
```


```
Total elements: 1000000
```

```
calculate_z_serial_purepython took 4.945113182067871 seconds
```

```
36221995 function calls in 5.263 seconds
```

```
Ordered by: cumulative time
```

ncalls	totttime	percall	cumtime	percall	filename:lineno(function)
1	0.000	0.000	5.263	5.263	{built-in method builtins.exec}
1	0.013	0.013	5.263	5.263	JuliaSet.py:1(<module>)
1	0.252	0.252	5.249	5.249	JuliaSet.py:21(calc_pure_python)
1	3.975	3.975	4.945	4.945	JuliaSet.py:59(calculate_z_serial_purepython)
34219980	0.970	0.000	0.970	0.000	{built-in method builtins.abs}
2002000	0.048	0.000	0.048	0.000	{method 'append' of 'list' objects}
1	0.004	0.004	0.004	0.004	{built-in method builtins.sum}
3	0.000	0.000	0.000	0.000	{built-in method builtins.print}
2	0.000	0.000	0.000	0.000	{built-in method time.time}
4	0.000	0.000	0.000	0.000	{built-in method builtins.len}
1	0.000	0.000	0.000	0.000	{method 'disable' of '_lsprof.Profiler' objects}



- Inside `calculate_z_serial_purepython`, the time spent on lines of code (without calling other functions) is 3.9 seconds.
- This function makes 34,219,980 calls to `abs`, which take a total of 0.9 seconds, along with other calls that do not cost much time.


```
stef@Stefs-MacBook-Air Codes % python -m cProfile -s cumulative JuliaSet.py
Length of x: 1000
Total elements: 1000000
calculate_z_serial_purepython took 4.945113182067871 seconds
      36221995 function calls in 5.263 seconds
```

Ordered by: cumulative time

ncalls	tottime	percall	cumtime	percall	filename:lineno(function)
1	0.000	0.000	5.263	5.263	{built-in method builtins.exec}
1	0.013	0.013	5.263	5.263	JuliaSet.py:1(<module>)
1	0.252	0.252	5.249	5.249	JuliaSet.py:21(calc_pure_python)
1	3.975	3.975	4.945	4.945	JuliaSet.py:59(calculate_z_serial_purepython)
34219980	0.970	0.000	0.970	0.000	{built-in method builtins.abs}
2002000	0.048	0.000	0.048	0.000	{method 'append' of 'list' objects}
1	0.004	0.004	0.004	0.004	{built-in method builtins.sum}
3	0.000	0.000	0.000	0.000	{built-in method builtins.print}
2	0.000	0.000	0.000	0.000	{built-in method time.time}
4	0.000	0.000	0.000	0.000	{built-in method builtins.len}
1	0.000	0.000	0.000	0.000	{method 'disable' of '_lsprof.Profiler' objects}

- While the per-call cost is negligible (it is recorded as 0.000 seconds), the total time for 34,219,980 calls is 0.97 seconds.
- We couldn't predict in advance exactly how many calls would be made to abs, as the Julia function has unpredictable dynamics

```

stef@Stefs-MacBook-Air Codes % python -m cProfile -s cumulative JuliaSet.py
Length of x: 1000
Total elements: 1000000
calculate_z_serial_purepython took 4.945113182067871 seconds
36221995 function calls in 5.263 seconds

```

Ordered by: cumulative time

ncalls	totttime	percall	cumtime	percall	filename:lineno(function)
1	0.000	0.000	5.263	5.263	{built-in method builtins.exec}
1	0.013	0.013	5.263	5.263	JuliaSet.py:1(<module>)
1	0.252	0.252	5.249	5.249	JuliaSet.py:21(calc_pure_python)
1	3.975	3.975	4.945	4.945	JuliaSet.py:59(calculate_z_serial_purepython)
34219980	0.970	0.000	0.970	0.000	{built-in method builtins.abs}
2002000	0.000	0.000	0.048	0.000	{method 'append' of 'list' objects}
1	0.004	0.004	0.004	0.004	{built-in method builtins.sum}
3	0.000	0.000	0.000	0.000	{built-in method builtins.print}
2	0.000	0.000	0.000	0.000	{built-in method time.time}
4	0.000	0.000	0.000	0.000	{built-in method builtins.len}
1	0.000	0.000	0.000	0.000	{method 'disable' of '_lsprof.Profiler' objects}

- The next line in the profiled output, {method 'append' of 'list' objects}, details the creation of 2,002,000 list items.
- This creation of 2,002,000 items is occurring in `calc_pure_python` during the setup phase.
- The `zs` and `cs` lists will be 1000*1000 items each (generating 1,000,000 * 2 calls), and these are built from a list of 1,000 `x` and 1,000 `y` coordinates. In total, this is 2,002,000 calls to append.

```
stef@Stefs-MacBook-Air Codes % python -m cProfile -s cumulative JuliaSet.py
Length of x: 1000
Total elements: 1000000
calculate_z_serial_purepython took 4.945113182067871 seconds
36221995 function calls in 5.263 seconds
```

Ordered by: cumulative time

ncalls	tottime	percall	cumtime	percall	filename:lineno(function)
1	0.000	0.000	5.263	5.263	{built-in method builtins.exec}
1	0.013	0.013	5.263	5.263	JuliaSet.py:1(<module>)
1	0.252	0.252	5.249	5.249	JuliaSet.py:21(calc_pure_python)
1	3.975	3.975	4.945	4.945	JuliaSet.py:59(calculate_z_serial_purepython)
34219980	0.970	0.000	0.970	0.000	{built-in method builtins.abs}
2002000	0.048	0.000	0.048	0.000	{method 'append' of 'list' objects}
1	0.004	0.004	0.004	0.004	{built-in method builtins.sum}
3	0.000	0.000	0.000	0.000	{built-in method builtins.print}
2	0.000	0.000	0.000	0.000	{built-in method time.time}
4	0.000	0.000	0.000	0.000	{built-in method builtins.len}
1	0.000	0.000	0.000	0.000	{method 'disable' of '_lsprof.Profiler' objects}



- The final line of the profiling output refers to `lsprof`
 - This is the original name of the tool that evolved into `cProfile` and can be ignored.



Write a Statistics File with cProfile

1 `python -m cProfile -o profile.stats JuliaSet.py` to obtain output file

```
stef@Stefs-MacBook-Air Codes % python
Python 3.8.9 (default, Aug 3 2021, 19:21:54)
[Clang 13.0.0 (clang-1300.0.29.3)] on darwin
Type "help", "copyright", "credits" or "license" for more information.
>>> import pstats
>>> p = pstats.Stats("profile.stats")
>>> p.sort_stats("cumulative")
<pstats.Stats object at 0x100439910>
>>> p.print_stats()
Sun Jan 16 15:41:03 2022      profile.stats
```



2 Python to read the output
`profile.stats`

36221995 function calls in 5.284 seconds

Ordered by: cumulative time

ncalls	tottime	percall	cumtime	percall	filename:lineno(function)
1	0.000	0.000	5.284	5.284	{built-in method builtins.exec}
1	0.014	0.014	5.284	5.284	JuliaSet.py:1(<module>)
1	0.245	0.245	5.270	5.270	JuliaSet.py:21(calc_pure_python)
1	3.983	3.983	4.973	4.973	JuliaSet.py:59(calculate_z_serial_purepython)
34219980	0.990	0.000	0.990	0.000	{built-in method builtins.abs}
2002000	0.048	0.000	0.048	0.000	{method 'append' of 'list' objects}
1	0.004	0.004	0.004	0.004	{built-in method builtins.sum}
3	0.000	0.000	0.000	0.000	{built-in method builtins.print}
2	0.000	0.000	0.000	0.000	{built-in method time.time}
4	0.000	0.000	0.000	0.000	{built-in method builtins.len}
1	0.000	0.000	0.000	0.000	{method 'disable' of '_lsprof.Profiler' objects}

<pstats.Stats object at 0x100439910>



Visualizing cProfile Output with SnakeViz

- `snakeviz` is a visualizer that draws the output of `cProfile` as a diagram in which larger boxes are areas of code that take longer to run.
- Use `snakeviz` to get a high-level understanding of a `cProfile` statistics file, particularly if you're investigating a new project for which you have little intuition.
 - The diagram will help you visualize the CPU-usage behavior of the system, and it may highlight areas that you hadn't expected to be expensive.
- To install SnakeViz, use `$ pip install snakeviz`
- To run `python -m snakeviz profile.stats --server`
- You will need the browser to open page given by `snakeviz`



Reset Root

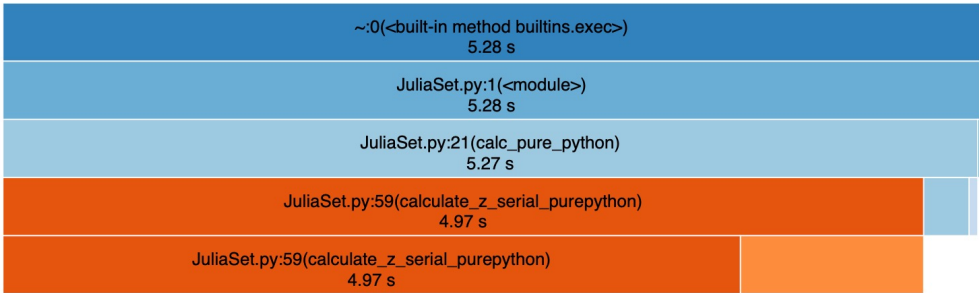
Reset Zoom

Style: **Icicle**

Depth: **10**

Cutoff: **1 / 1000**

Call Stack



ncalls	tottime	percall	cumtime	percall	filename:lineno(function)
1	3.983	3.983	4.973	4.973	JuliaSet.py:59(calculate_z_serial_purepython)
34219980	0.9902	2.894e-08	0.9902	2.894e-08	~:0(<built-in method builtins.abs>)
1	0.2446	0.2446	5.27	5.27	JuliaSet.py:21(calc_pure_python)
2002000	0.04771	2.383e-08	0.04771	2.383e-08	~:0(<method 'append' of 'list' objects>)
1	0.01426	0.01426	5.284	5.284	JuliaSet.py:1(<module>)
1	0.004439	0.004439	0.004439	0.004439	~:0(<built-in method builtins.sum>)
3	9.175e-05	3.058e-05	9.175e-05	3.058e-05	~:0(<built-in method builtins.print>)
2	2.875e-06	1.438e-06	2.875e-06	1.438e-06	~:0(<built-in method time.time>)
1	1.417e-06	1.417e-06	5.284	5.284	~:0(<built-in method builtins.exec>)
4	5.84e-07	1.46e-07	5.84e-07	1.46e-07	~:0(<built-in method builtins.len>)
1	3.33e-07	3.33e-07	3.33e-07	3.33e-07	~:0(<method 'disable' of '_Isprof.Profiler' objects>)

To Summarize

- `cProfile` is a built-in Python profiler to measure the time taken to run every function.
- `cProfile` is ideal to find the most computational-intensive parts of the code.
- It is typically called from the command line with `-m cProfile` and can print to stdout or to a file
- The output file can be visualized with the `SnakeViz` tool.