




DD2358 – Using `line_profiler` for Line-by-Line Measurements

Stefano Markidis

KTH Royal Institute of Technology





Intended Learning Outcomes


- To deploy the `line_profiler` module in your code to profile it line-by-line
- Analyze the `line_profiler` statistic output and identify fine-grained computational bottle-necks

line_profiler

- `line_profiler` works by profiling individual functions on a line-by-line basis
- When profiling, we should start with `cProfile` and use the high-level view to guide which functions to profile with `line_profiler`.
- To use `line_profiler`
 1. You need to install `line_profiler` with `pip install line_profiler`
 2. A decorator (`@profile`) is used to **mark the chosen function**.
 3. The **`kernprof` script** is used to execute our code
 - the CPU time and other statistics for each line of the chosen function are recorded.

Add @profile Decorator before Function

@profile



```
def calculate_z_serial_purepython(maxiter, zs, cs):  
    """Calculate output list using Julia update rule"""  
    output = [0] * len(zs)  
    for i in range(len(zs)):  
        n = 0  
        z = zs[i]  
        c = cs[i]  
        while abs(z) < 2 and n < maxiter:  
            z = z * z + c  
            n += 1  
        output[i] = n  
    return output
```



Run kernprof for Obtaining Profiling

```
python -m kernprof -l JuliaSet.py
```

Length of x: 1000

Total elements: 1000000


calculate_z_serial_purepython took 54.3584349155426 seconds

Wrote profile results to `JuliaSet.py.lprof`

Note the timing!



Profiling information
Stored in `.lprof`



Print profile: `python -m line_profiler JuliaSet.py.lprof`

```
stef@Stefs-MacBook-Air Codes % python -m line_profiler JuliaSet.py.lprof
```

```
Timer unit: 1e-06 s
```

```
Total time: 29.5089 s
```

```
File: JuliaSet.py
```

```
Function: calculate_z_serial_purepython at line 58
```

Line #	Hits	Time	Per Hit	% Time	Line Contents
58					@profile
59					def calculate_z_serial_purepython(maxiter, zs, cs):
60					"""Calculate output list using Julia update rule"""
61	1	1562.0	1562.0	0.0	output = [0] * len(zs)
62	1000001	241283.0	0.2	0.8	for i in range(len(zs)):
63	1000000	228023.0	0.2	0.8	n = 0
64	1000000	273051.0	0.3	0.9	z = zs[i]
65	1000000	302535.0	0.3	1.0	c = cs[i]
66	34219980	10643253.0	0.3	36.1	while abs(z) < 2 and n < maxiter:
67	33219980	9110298.0	0.3	30.9	z = z * z + c
68	33219980	8445538.0	0.3	28.6	n += 1
69	1000000	263390.0	0.3	0.9	output[i] = n
70	1	5.0	5.0	0.0	return output

```
stef@Stefs-MacBook-Air Codes % python -m line_profiler JuliaSet.py.lprof
Timer unit: 1e-06 s
```

Total time: 29.5089 s

File: JuliaSet.py

Function: calculate_z_serial_purepython at line 58

Line #	Hits	Time	Per Hit	% Time	Line Contents
58					@profile
59					def calculate_z_serial_purepython(maxiter, zs, cs):
60					"""Calculate output list using Julia update rule"""
61	1	1562.0	1562.0	0.0	output = [0] * len(zs)
62	1000001	241283.0	0.2	0.8	for i in range(len(zs)):
63	1000000	228023.0	0.2	0.8	n = 0
64	1000000	273051.0	0.3	0.9	z = zs[i]
65	1000000	302535.0	0.3	1.0	c = cs[i]
66	34219980	10643253.0	0.3	36.1	while abs(z) < 2 and n < maxiter:
67	33219980	9110298.0	0.3	30.9	z = z * z + c
68	33219980	8445538.0	0.3	28.6	n += 1
69	1000000	263390.0	0.3	0.9	output[i] = n
70	1	5.0	5.0	0.0	return output

- The **% Time column** is the most helpful - we can see that 36% of the time is spent on the while testing.
 - We don't know whether the first statement ($\text{abs}(z) < 2$) is more expensive than the second ($n < \text{maxiter}$), though.
 - Inside the loop, we see that the update to z is also fairly expensive: 30.
- **Even $n += 1$ is expensive!** Python's dynamic lookup machinery is at work for every loop, even though we're using the same types for each variable in each loop
- The creation of the outputlist and the updates on line 20 are relatively cheap compared to the cost of the while loop.

Performance Improvement Opportunities

```
stef@Stefs-MacBook-Air Codes % python -m line_profiler JuliaSet.py.lprof
Timer unit: 1e-06 s
```

Total time: 29.5089 s

File: JuliaSet.py

Function: calculate_z_serial_purepython at line 58

Line #	Hits	Time	Per Hit	% Time	Line Contents
=====					
58					@profile
59					def calculate_z_serial_purepython(maxiter, zs, cs):
60					"""Calculate output list using Julia update rule"""
61	1	1562.0	1562.0	0.0	output = [0] * len(zs)
62	1000001	241283.0	0.2	0.8	for i in range(len(zs)):
63	1000000	228023.0	0.2	0.8	n = 0
64	1000000	273051.0	0.3	0.9	z = zs[i]
65	1000000	302535.0	0.3	1.0	c = cs[i]
66	34219980	10643253.0	0.3	36.1	while abs(z) < 2 and n < maxiter:
67	33219980	9110298.0	0.3	30.9	z = z * z + c
68	33219980	8445538.0	0.3	28.6	n += 1
69	1000000	263390.0	0.3	0.9	output[i] = n
70	1	5.0	5.0	0.0	return output



To Summarize

- The `line_profiler` module allows us to profile individual functions on a line-by-line basis using the `@profile` decorator in your code.
- The analysis of the profile information (the `.lsprof` file) shows us code lines with potential to be optimized.